

Concurrency and Parallel Programming

Assignment 2

David van Erkelens and Jelte Fennema
Department of Computer Science
University of Amsterdam


November 15, 2012

1 Notes

- The smallest problem size doesn't benefit from MPI in any of the cases. This was also the case with threading.
- For some reason the tests with problem size 10^5 performed abnormally bad in all of the tests. It only benefits a bit from MPI. Way less than any of the others. (except for 10^3) Both 10^4 and 10^6 perform way better in all cases.
- With the two largest problem sizes MPI gains around the maximum theoretical speedup when running on two different nodes with one process per node. In comparison with using threading MPI performs very well in these cases.
- As seen in the graphs, the speedup depends on the problem size. When the problem size is small, the speedup is also very small and sometimes even smaller than 1. With a large problem size, the speedup can reach almost 8. This can be explained by the fact that there can be a large communication delay with a small problem size. However, when there is a larger problem size (and thus more calculations), the communication delay will be a smaller percentage of the total executing time.

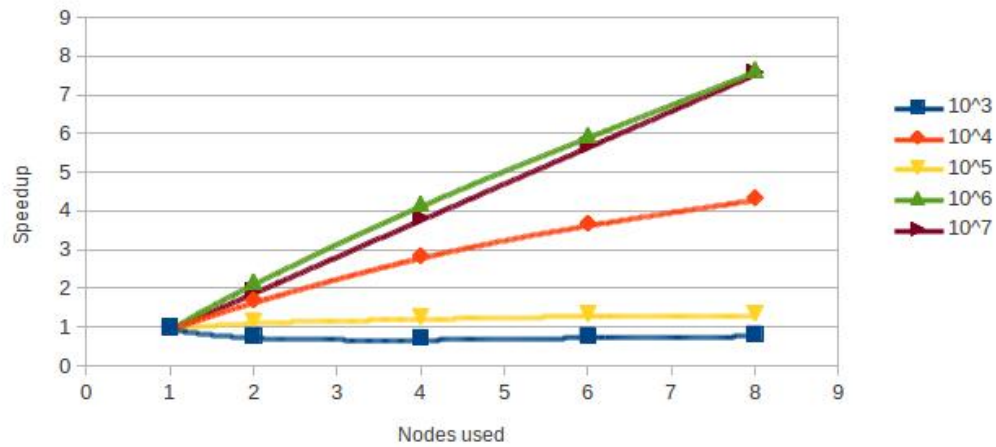
- As seen in the graphs, all problem sizes quickly reach their maximum possible speedup. Only the largest problem size doesn't stagnate yet. However, when the problem will be executed on even more cores, this problem size will probably also stagnate, but there's no way to retrieve this from the current results.

2 Graphs



Problem size vs nodes

With one process per node



Problem size vs nodes

With eight processes per node

